

import the libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import sklearn

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn import metrics
from sklearn.metrics import accuracy_score
from sklearn import svm
from sklearn.tree import DecisionTreeClassifier
```

load the data

```
In [2]: data = pd.read_csv("/home/tamanna/Downloads/creditcard.csv")
data
```

Out[2]:

	Time	V1	V2	V3	V4	V5	V6	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.
...
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.0
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.

284807 rows × 31 columns

data preprocessing

```
In [3]: data.head()
```

Out[3]:

	Time	V1	V2	V3	V4	V5	V6	V7	
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.0
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.0
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.2
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.3
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.2

5 rows × 31 columns

In [4]: data.tail()

Out[4]:

	Time	V1	V2	V3	V4	V5	V6	V7	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.4	
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.0	
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.0	
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.0	
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.5	

5 rows × 31 columns

In [5]: # generating descriptive statistics
data.describe()

Out[5]:

	Time	V1	V2	V3	V4	V5	V6	V7
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604095e-16	-1.137171e-15	-1.137171e-15
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380171e+00	1.380171e+00	1.380171e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137171e-15	-1.137171e-15	-1.137171e-15
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.911711e-01	-6.911711e-01	-6.911711e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433171e-01	-5.433171e-01	-5.433171e-01
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119171e-01	6.119171e-01	6.119171e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480171e+00	3.480171e+00	3.480171e+00

8 rows × 31 columns

In [6]: data.shape

Out[6]: (284807, 31)

```
In [7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        284807 non-null  float64
1   V1          284807 non-null  float64
2   V2          284807 non-null  float64
3   V3          284807 non-null  float64
4   V4          284807 non-null  float64
5   V5          284807 non-null  float64
6   V6          284807 non-null  float64
7   V7          284807 non-null  float64
8   V8          284807 non-null  float64
9   V9          284807 non-null  float64
10  V10         284807 non-null  float64
11  V11         284807 non-null  float64
12  V12         284807 non-null  float64
13  V13         284807 non-null  float64
14  V14         284807 non-null  float64
15  V15         284807 non-null  float64
16  V16         284807 non-null  float64
17  V17         284807 non-null  float64
18  V18         284807 non-null  float64
19  V19         284807 non-null  float64
20  V20         284807 non-null  float64
21  V21         284807 non-null  float64
22  V22         284807 non-null  float64
23  V23         284807 non-null  float64
24  V24         284807 non-null  float64
25  V25         284807 non-null  float64
26  V26         284807 non-null  float64
27  V27         284807 non-null  float64
28  V28         284807 non-null  float64
29  Amount      284807 non-null  float64
30  Class       284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

```
In [8]: data.isnull().sum()
```


```
Out[8]: Time      0
        V1        0
        V2        0
        V3        0
        V4        0
        V5        0
        V6        0
        V7        0
        V8        0
        V9        0
        V10       0
        V11       0
        V12       0
        V13       0
        V14       0
        V15       0
        V16       0
        V17       0
        V18       0
        V19       0
        V20       0
        V21       0
        V22       0
        V23       0
        V24       0
        V25       0
        V26       0
        V27       0
        V28       0
        Amount    0
        Class     0
        dtype: int64
```

```
In [9]: data['Class'].value_counts()
```

```
Out[9]: Class
0      284315
1         492
Name: count, dtype: int64
```

```
In [10]: data['Class'].value_counts().plot(kind = 'bar')
```

```
Out[10]: <AxesSubplot: xlabel='Class'>
```



dataset is highly imbalanced.

```
In [11]: non_fraud = data[data.Class == 0]
         fraud = data[data.Class == 1]
```

```
In [12]: non_fraud.shape
```

```
Out[12]: (284315, 31)
```

```
In [13]: fraud.shape
```

```
Out[13]: (492, 31)
```

```
In [14]: non_fraud_sample = non_fraud.sample(600)
```

```
In [15]: non_fraud_sample
```

Out[15]:

	Time	V1	V2	V3	V4	V5	V6	
79938	58251.0	-1.053308	0.588041	0.497243	0.514123	-0.243353	-1.143231	0.5
94158	64757.0	-4.283375	3.543912	-2.966651	-1.359581	0.587368	2.866529	-1.74
49211	43957.0	1.018221	0.085412	-0.353711	1.066409	0.340708	-0.150033	0.48
267573	162831.0	2.133450	0.005781	-2.323607	0.049250	0.971403	-0.498335	0.5
138673	82787.0	-0.935118	1.202229	1.257032	0.404653	0.133229	0.160176	0.10
...	
212291	138799.0	-0.878858	0.645186	-0.272966	-1.193022	3.076594	3.633350	0.2
27185	34430.0	1.209556	-0.166461	1.181014	0.565935	-0.981248	-0.130360	-0.78
93551	64483.0	1.105351	-1.457671	0.749924	-2.015344	-1.728027	-0.056765	-1.1
185371	126638.0	2.101841	0.625172	-2.668427	0.594459	0.963007	-1.364090	0.4
283626	171743.0	0.425363	-0.062341	-0.857997	0.326586	1.822954	-0.914518	0.30

600 rows × 31 columns

In [16]:

new_data = pd.concat([non_fraud_sample, fraud], axis =0)

In [17]:

new_data

Out[17]:

	Time	V1	V2	V3	V4	V5	V6	
79938	58251.0	-1.053308	0.588041	0.497243	0.514123	-0.243353	-1.143231	0.51
94158	64757.0	-4.283375	3.543912	-2.966651	-1.359581	0.587368	2.866529	-1.74
49211	43957.0	1.018221	0.085412	-0.353711	1.066409	0.340708	-0.150033	0.48
267573	162831.0	2.133450	0.005781	-2.323607	0.049250	0.971403	-0.498335	0.51
138673	82787.0	-0.935118	1.202229	1.257032	0.404653	0.133229	0.160176	0.10
...	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.88
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.41
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.23
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.20
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.22

1092 rows × 31 columns

values of x and y

```
In [18]: x = new_data.drop(['Class'], axis = 1)
```

```
In [19]: print(x)
```

	Time	V1	V2	V3	V4	V5	V6	
\								
79938	58251.0	-1.053308	0.588041	0.497243	0.514123	-0.243353	-1.143231	
94158	64757.0	-4.283375	3.543912	-2.966651	-1.359581	0.587368	2.866529	
49211	43957.0	1.018221	0.085412	-0.353711	1.066409	0.340708	-0.150033	
267573	162831.0	2.133450	0.005781	-2.323607	0.049250	0.971403	-0.498335	
138673	82787.0	-0.935118	1.202229	1.257032	0.404653	0.133229	0.160176	
...	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	
	V7	V8	V9	...	V20	V21	V22	\
79938	0.516954	0.188210	-0.619711	...	-0.083019	0.128897	0.083167	
94158	-1.748133	3.303362	-0.647340	...	0.147722	0.057162	-0.549032	
49211	0.484130	-0.056734	-0.563302	...	0.021810	0.117429	0.155288	
267573	0.518368	-0.259087	-0.017044	...	-0.206715	0.098535	0.411672	
138673	0.103757	0.633826	-1.232453	...	0.066806	-0.115048	-0.531792	
...	
279863	-0.882850	0.697211	-2.064945	...	1.252967	0.778584	-0.319189	
280143	-1.413170	0.248525	-1.127396	...	0.226138	0.370612	0.028234	
280149	-2.234739	1.210158	-0.652250	...	0.247968	0.751826	0.834108	
281144	-2.208002	1.058733	-1.632333	...	0.306271	0.583276	-0.269209	
281674	0.223050	-0.068384	0.577829	...	-0.017652	-0.164350	-0.295135	
	V23	V24	V25	V26	V27	V28	Amount	
79938	0.014889	0.415131	-0.413715	0.457807	-0.151270	0.062739	72.98	
94158	0.024283	1.006805	1.059677	0.419387	-0.541254	0.052017	6.76	
49211	-0.250932	-0.289248	0.743356	-0.257872	-0.017277	0.010981	100.75	
267573	-0.081135	0.165554	0.441014	0.700315	-0.124472	-0.095485	0.76	
138673	-0.110064	-0.339523	-0.035581	0.308557	-0.068997	-0.007041	9.99	
...	
279863	0.639419	-0.294885	0.537503	0.788395	0.292680	0.147968	390.00	
280143	-0.145640	-0.081049	0.521875	0.739467	0.389152	0.186637	0.76	
280149	0.190944	0.032070	-0.739695	0.471111	0.385107	0.194361	77.89	
281144	-0.456108	-0.183659	-0.328168	0.606116	0.884876	-0.253700	245.00	
281674	-0.072173	-0.450261	0.313267	-0.289617	0.002988	-0.015309	42.53	

[1092 rows x 30 columns]

```
In [20]: y = new_data["Class"]
```

```
In [21]: print(y)
```

```

79938      0
94158      0
49211      0
267573     0
138673     0
..
279863     1
280143     1
280149     1
281144     1
281674     1
Name: Class, Length: 1092, dtype: int64

```

splitting of data into training and testing sets

```
In [22]: x_train, x_test, y_train, y_test = train_test_split(x,y, random_state = 2, t
```

```
In [23]: print(x.shape, x_train.shape, x_test.shape)
```

```
(1092, 30) (982, 30) (110, 30)
```

Logistic Regression

```
In [24]: model = LogisticRegression()
```

```
In [25]: model.fit(x_train, y_train)
```

```
Out[25]: ▼ LogisticRegression
LogisticRegression()
```

Model Evaluation

```
In [26]: x_train_prediction = model.predict(x_train)
training_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
In [27]: training_data_accuracy
```

```
Out[27]: 0.9480651731160896
```

```
In [28]: x_test_prediction = model.predict(x_test)
testing_data_accuracy = accuracy_score(x_test_prediction, y_test)
```

```
In [29]: testing_data_accuracy
```

```
Out[29]: 0.9454545454545454
```

svm classifier

```
In [30]: model2 = svm.SVC(kernel = "linear")
```



```
In [31]: model2.fit(x_train, y_train)
```

```
Out[31]: SVC
SVC(kernel='linear')
```

```
In [32]: x_train_prediction = model2.predict(x_train)
```

Model Evaluation

```
In [33]: training_data_accuracy = accuracy_score(x_train_prediction, y_train)
```

```
In [34]: training_data_accuracy
```

```
Out[34]: 0.8981670061099797
```

```
In [35]: x_test_prediction = model2.predict(x_test)
```

```
In [36]: testing_data_accuracy = accuracy_score(x_test_prediction, y_test)
testing_data_accuracy
```

```
Out[36]: 0.9181818181818182
```

Decision Tree Classifier

```
In [37]: model3 = DecisionTreeClassifier(max_leaf_nodes = 10, random_state = 0)
model3.fit(x_train, y_train)
```

```
Out[37]: DecisionTreeClassifier
DecisionTreeClassifier(max_leaf_nodes=10, random_state=0)
```

```
In [39]: x_train_prediction = model3.predict(x_train)
```

Model Evaluation

```
In [40]: training_data_accuracy = accuracy_score(x_train_prediction, y_train)
training_data_accuracy
```

```
Out[40]: 0.955193482688391
```

```
In [41]: x_test_prediction = model3.predict(x_test)
```

```
In [42]: testing_data_accuracy = accuracy_score(x_test_prediction, y_test)
```

```
In [43]: testing_data_accuracy
```

```
Out[43]: 0.9181818181818182
```

logistic regression has got the highest accuracy as compared to both decision tree classifier and svm